Milestone 1 Review Meeting Summary

HeartCompass / Digital Immortality

Project	HeartCompass / Digital Immortality
Meeting Time	May 9, 2026, 4:00 PM
Duration	44 minutes 42 seconds
Presenter	Tian Bu
Prepared By	Leishiyu Chen, Haojie Hu, Tian Bu, Yihan Liu
Document Date	May 9, 2026
Submission Purpose	Canvas submission for Milestone 1 Review Meeting Summary

1. Main Content of the Meeting

The presenter introduced Digital Immortality as a persona-centric AI project built to solve the consumption problem left by file-based persona-distillation skills. Existing skills can extract personality, interaction, procedural, and memory materials, but they tend to grow into large files that must be read in full or nearly in full. The project therefore aims to engineer persona construction and persona consumption into a maintainable system.

The team explained the current product form: a CLI-first management tool and a Lark Bot for daily use. The team intentionally does not prioritize a traditional web front end at this stage, because the main challenge is the AI workflow, memory structure, data layer, and low-friction interaction through an IM channel.

The core architecture was presented as layered storage plus vectorized recall. Lightweight intrinsic fields are stored directly in FigureAndRelation, while large, high-context persona facts are stored as FineGrainedFeed records by dimension and retrieved semantically according to the current conversation context.

Two main graph workflows were discussed. FRBuildingGraph receives raw user materials, cleans and preprocesses them, extracts intrinsic fields and fine-grained feeds, compares new values against stored values, records conflicts, and writes update logs. ConversationGraph receives batched Lark messages, loads persona context, recalls relevant feeds, manages short-term memory, trims and summarizes older rounds, and generates persona-grounded replies.

The meeting also covered the rationale for LangGraph rather than a pure ReAct agent. The team emphasized that persona construction and dialogue generation require controllable, observable, and repeatable steps instead of a black-box autonomous agent.

The team introduced Harness work as a way to convert repeated engineering practices into reusable skills and rules. Examples included code-quality review, commit-update writing, document generation, document optimization, graph documentation, and session-practice learning.

2. Reviewer Comments and Advice

Reviewer	Focus	Comment / Advice	Team Response
Ding Zhenyao	Information retention and iteration	Asked whether the system can preserve original persona information during later iterations and avoid losing foundational facts.	The team explained that stored persona data will remain in the database, while corrections or missing information can be handled through the persona-building update flow.
Chen Yiming	Differentiation from existing persona skills	Asked how the project differs from existing colleague/ex-partner/persona-distillation skills.	The team clarified that existing skills mainly produce files for another agent to consume, while this project provides an end-to-end engineered product with database storage, vector recall, CLI/Lark access, and no dependency on a general-purpose external agent workflow.
Li Ranxi	Project management artifacts	Praised the project architecture and delivery thinking, but noted that the presentation spent more time on technical design than on RBS, WBS, planned hours, and schedule progress.	The team accepted this as a key improvement area and updated the RBS, WBS, schedule, and task-hour allocation for the next submission.
Huang Yihe	Persona consistency and memory realism	Suggested using an LLM to periodically review indirect logical conflicts or self-contradictions in persona items. Also suggested making memory more realistic by considering time decay and event importance.	The team identified this as a next-milestone action item for persona health review, conflict inspection, and memory-decay design.
Qian Baoqiang	Message buffering responsiveness	Pointed out that the 15-second Lark message buffer may reduce responsiveness, because users may want a timely reply even while sending a series of short messages.	The team will evaluate adaptive batching, early response triggers, and configurable waiting time instead of relying only on a fixed 15-second reset window.
Shangguan Siyang	Ethical boundary and abuse prevention	Asked whether the system has usage boundaries for simulating real people and how misuse can be prevented.	The team discussed role-based extraction limits, especially stricter boundaries for colleagues, mentors, and public figures, and the need to restrict full persona export after the system becomes more mature.
Zhao Zhuobing	Next-step roadmap	Asked about the next planned work after the current milestone.	The team listed single-Lark-bot limitations, shared database support, lower setup barriers, built-in public figures, prompt optimization, Harness work, new correction/analysis/decision graphs, log governance, and external long-term memory integration.

3. Action Plan for Current Problems in the Next Milestone

The next milestone will focus on converting reviewer feedback into concrete management and engineering tasks. The first action has already been reflected in the updated RBS, WBS, schedule, and planned working hours below.

ID	Action	Description	Owner	Target	Status
A1	Update RBS, WBS, schedule, and task-hour allocation	Revise project-management artifacts to show original and updated scope, task-level decomposition, planned hours, and total hours. This action aligns with WP5.6 final documentation and governance assets.	Tian Bu (team leader)	W11 / before Canvas submission	Completed
A2	Add persona consistency review	Design a periodic LLM-based review flow to detect indirect conflicts, contradictions, stale facts, and unresolved FineGrainedFeedConflict records. This action aligns with WP2.3 conflict detection and WP2.5 persona-building review.	Leishiyu Chen	W6 / Milestone 2	Planned
A3	Improve memory realism	Introduce memory-decay thinking, event importance, and time-aware recall behavior so long-term memories can feel more realistic. This action aligns with WP3.1 service-backed recall and WP3.2 short-term memory trim and summary.	Haojie Hu	W6-W7 / Milestone 2	Planned
A4	Optimize Lark message buffering	Evaluate adaptive batching, shorter waits, manual reply triggers, and configuration options for WAITING_SECONDS_FOR_CONVERSATION. This action aligns with WP4.1 Lark WebSocket handling and WP4.3 message batching.	Yihan Liu	W6-W8 / Milestone 2	Planned
A5	Define ethical and role-based boundaries	Document what can be extracted, recalled, displayed, or exported for self, family, friends, colleagues, mentors, public figures, and strangers. This action aligns with WP1.3 FigureRole management and WP4.2 bot menu permissions.	Tian Bu (team leader)	W4-W7 / Milestone 2	Planned
A6	Lower setup barriers with shared database design	Investigate encrypted shared database support, gateway-based access, and protection against exposing database URIs to users. This action aligns with WP5.3 Docker/database automation and WP5.4 regression testing.	Leishiyu Chen	W9-W10 / Milestone 2-Milestone 3	Planned
A7	Address single Lark Bot limitation	Design multi-bot or multi-service support so a single installed package is not limited to one Lark Bot proxy. This action aligns with WP4.4 service-only data access alignment and WP4.5 channel integration testing.	Yihan Liu	W8 / Milestone 2	Planned
A8	Strengthen log governance	Reduce heavy ConversationGraph logging, preserve necessary checkpoint observability, and avoid latency from excessive log payloads. This action aligns with WP5.1 CLI logs and WP5.4 graph/channel/CLI regression testing.	Haojie Hu	W8-W10 / Milestone 2	Planned

4. Original and Updated RBS

The original RBS described five broad responsibility streams. After the review, the RBS was updated into a numbered hierarchy that better matches the current main-branch implementation and gives traceability to the updated WBS.

Original RBS

ID	Original Responsibility Area	Scope
R1	User and Relationship Data Foundation	User registration, profile management, MBTI collection, Crush and RelationChain setup, ownership and validation rules.
R2	Context Acquisition and Knowledge Management	Natural language narratives, screenshots, curated knowledge, Event, ChatTopic, DerivedInsight, Knowledge, and ContextEmbedding records.
R3	Relationship Analysis Engine	Communication suggestions, risk prompts, narrative analysis, conversation analysis, and interaction signals.
R4	Virtual Figure Interaction System	Persona context loading, memory recall, WebSocket interaction, timed message batching, and context refresh.
R5	Client Experience, Integration, and Deployment	Lightweight client, product integration, testing, performance optimization, packaging, and final delivery.

Updated RBS

ID	Updated Responsibility Area	Scope
R1.1	User and Figure-Relation Data Foundation	User identity, authentication, profile management, MBTI collection, FigureAndRelation records, FigureRole assignment, and ownership validation.
R1.2	Persona Building and Knowledge Management	FRBuildingGraph preprocessing, OriginalSource persistence, FineGrainedFeed extraction, conflict tracking, update logs, embeddings, and semantic recall.
R2.1	Conversation and Dialogue System	ConversationGraph persona loading, service-backed recall, prompt assembly, response generation, short-term memory trimming, rolling summary, and feed sync.
R2.2	Channel Integration and Lark Bot	Lark WebSocket access, open_id mapping, timed message batching, bot commands, card output, FR switching, and busy-state feedback.
R3.1	CLI, Harness, Packaging, and Deployment	immortality CLI, doctor/setup/auth/fr/logs/lark-service commands, JSON contracts, Docker/Alembic automation, PyPI packaging, Harness skills, testing, and deployment.

Updated RBS Chart

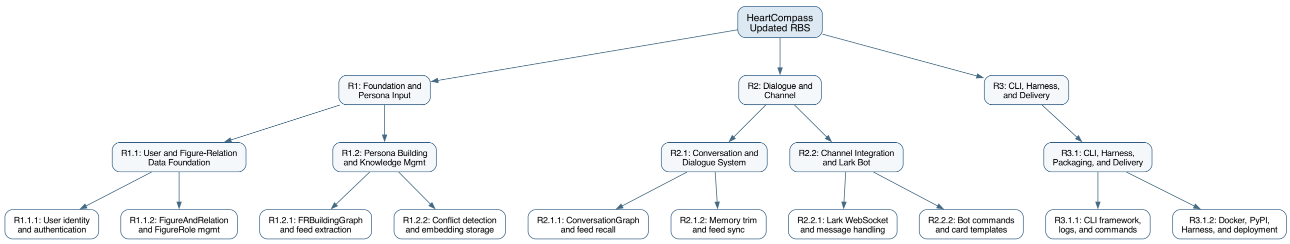

Figure 1. Updated Responsibility Breakdown Structure with numbered hierarchy.

5. Original WBS and Planned Working Hours

The original WBS used five major work packages with package-level planned hours. It did not yet provide sub-task-level decomposition, which was one of the most important project-management issues raised during the review.

WP	Original Work Package	Original Scope	Original Schedule	Hours
WP1	User and Relationship Data Foundation	Data model and user/relationship baseline	Mar 30-Apr 26	84
WP2	Context Acquisition and Knowledge Management	Context ingestion, screenshot/topic extraction, knowledge and embeddings	Apr 13-Apr 26	120
WP3	Relationship Analysis Engine	Narrative analysis, conversation analysis, scoring and guidance	Apr 20-May 10	132
WP4	Virtual Figure Interaction System	Persona loading, memory recall, WebSocket interaction and virtual figure dialogue	Apr 27-May 17	126
WP5	Client Experience, Integration, and Deployment	Lightweight client, integration, testing, optimization, and deployment	Apr 27-Jun 14	162

WP	Original Work Package	Original Scope	Original Schedule	Hours
Total				624

6. Updated WBS, Planned Working Hours, and Total Hours

The updated WBS preserves the same total planned workload of 624 hours, but it replaces the broad original work packages with traceable task IDs, RBS references, week assignments, and planned hours for each task.

WP	Updated Work Package	RBS	Hours
WP1	User and Figure-Relation Data Foundation	R1.1	84
WP2	Persona Building and Knowledge Management	R1.2	120
WP3	Conversation and Dialogue System	R2.1	132
WP4	Channel Integration and Lark Bot	R2.2	126
WP5	CLI, Harness, Packaging, and Deployment	R3.1	162
Total			624

Detailed Updated Task-Level WBS

Task ID	Task Name	RBS	Weeks	Hours
WP1.1	User registration and authentication	R1.1	W3	20
WP1.2	Profile management and MBTI collection	R1.1	W3	22
WP1.3	FigureAndRelation CRUD and FigureRole	R1.1	W4	24
WP1.4	Data foundation review and test	R1.1	W4	18
WP2.1	FRBuildingGraph preprocessing pipeline	R1.2	W4-W5	28
WP2.2	Fine-grained feed extraction and storage	R1.2	W5	26
WP2.3	Conflict detection and update logging	R1.2	W5-W6	22
WP2.4	Embedding generation and vector recall	R1.2	W6	26
WP2.5	Persona building review and test	R1.2	W6	18
WP3.1	Service-backed ConversationGraph recall	R2.1	W5-W6	28
WP3.2	Short-term memory trim and summary	R2.1	W6-W7	30
WP3.3	LLM call and response generation	R2.1	W7	28
WP3.4	Feed sync and prompt alignment	R2.1	W7-W8	26
WP3.5	Conversation system review and test	R2.1	W8	20
WP4.1	Lark WebSocket and message handling	R2.2	W6-W7	30
WP4.2	Bot menu and persona build commands	R2.2	W7	26
WP4.3	Message batching and card templates	R2.2	W7-W8	24
WP4.4	Service-only data access alignment	R2.2	W8	26
WP4.5	Busy-state cards and integration test	R2.2	W8	20
WP5.1	CLI framework, doctor/setup, and logs	R3.1	W7-W8	32
WP5.2	CLI auth, FR management, and JSON contracts	R3.1	W8-W9	30
WP5.3	Docker support and database automation	R3.1	W9	30
WP5.4	Graph/channel/CLI regression testing	R3.1	W9-W10	30
WP5.5	PyPI packaging and publishing	R3.1	W10	22
WP5.6	Final docs, Harness assets, and deployment	R3.1	W11	18
Total				624

Updated WBS Chart

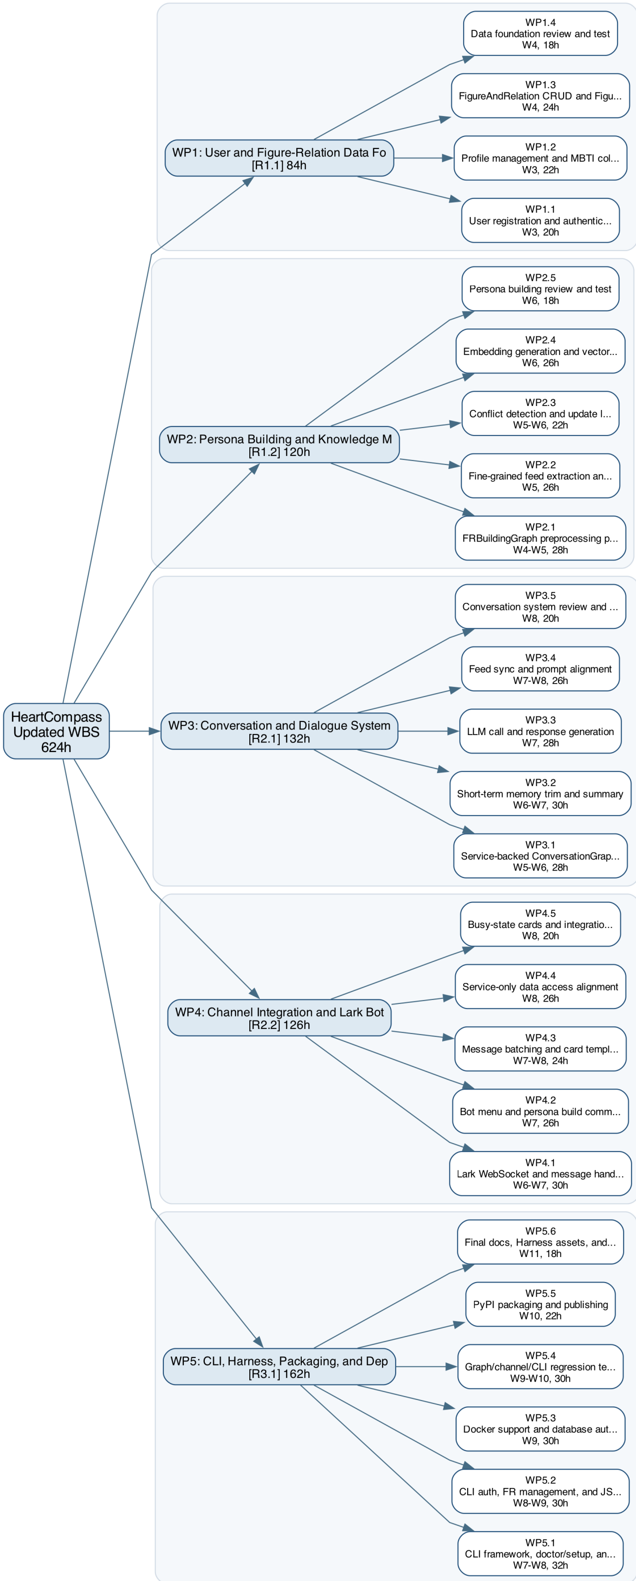

Figure 2. Updated Work Breakdown Structure with task-level work packages and planned hours.

7. Updated Schedule for the Assignment 'Update RBS, WBS and Schedule'

The schedule was updated from broad phase-level planning to week-by-week task completion. The new schedule allows the team to report concrete task IDs at weekly checkpoints and directly connects progress tracking to the WBS.

Week	Dates	Planned Completion / Milestone
W1	Mar 30-Apr 05	P0: Planning and Requirements
W2	Apr 06-Apr 12	P1: System Design
W3	Apr 13-Apr 19	WP1.1, WP1.2
W4	Apr 20-Apr 26	WP1.3, WP1.4, WP2.1 begins
W5	Apr 27-May 03	WP2.1, WP2.2, WP3.1 begins
W6	May 04-May 10	WP2.3, WP2.4, WP2.5, WP3.1, WP4.1 begins
W7	May 11-May 17	WP3.2, WP3.3, WP4.1, WP4.2
W8	May 18-May 24	WP3.4, WP3.5, WP4.3, WP4.4, WP4.5, WP5.1
W9	May 25-May 31	WP5.2, WP5.3
W10	Jun 01-Jun 07	WP5.4, WP5.5
W11	Jun 08-Jun 14	WP5.6

8. Summary of Changes from Original Plan to Updated Plan

- The project scope moved away from relationship-analysis APIs and a lightweight client as the primary delivery surface, and toward persona construction, ConversationGraph dialogue, CLI management, Lark Bot interaction, and Harness assets.

- Original entities such as Crush, RelationChain, Event, ChatTopic, DerivedInsight, and ContextEmbedding were replaced or narrowed into the current main-branch model centered on User, FigureAndRelation, OriginalSource, FineGrainedFeed, FineGrainedFeedConflict, FROverallUpdateLog, FRBuildingGraphReport, Knowledge, and Analysis.

- The updated RBS is numbered and traceable, while the original RBS was descriptive but not task-linked.

- The updated WBS keeps the same total 624 planned hours but introduces 25 task-level work items, week assignments, RBS references, and explicit review/test tasks.

- The updated schedule reflects current implementation realities: service-only data access, FRBuildingGraph concurrency control, CLI logs, Lark cards and commands, PyPI packaging, Docker/Alembic deployment support, and Harness documentation.

9. Reviewer Signatures

The following signatures confirm that the reviewers participated in the Milestone 1 review discussion and provided feedback recorded in this meeting summary.

Ding Zhenyao 丁桢垚	**Chen Yiming** 陈奕名
Li Ranxi 李冉[illegible]	**Huang Yihe** 黄一和
Qian Baoqiang 钱宝强	**Shangguan Siyang** 上官思洋
Zhao Zhuobing 赵卓冰	